

SQL for Analytics



SATYAJIT PATTNAIK

SQL Timeline




SQL Timeline Contd..



Database

Organised collection of Data stored in an electronic format




Database Management System

A System Software for creating and managing database



Database Architecture



Database Architecture



File Server

- ✓ Client request a file from the disk or a file-server
- ✓ Software opens the files, stores it in memory, makes changes in memory, and then saves it back to the disk drive (or file server)
- ✓ What happens if someone else opens the same file after you do, then you save your changes, and then they save their changes?
 - Potential data loss
 - Solution is file-locking
 - File locking is very inconvenient however as it means that nobody else can edit the file if someone else has it open first. (Logically good for small companies)

Client Server

- ✓ Desktop software requests data from the database
- ✓ Data is held in memory but when a change is made, a request (SQL statement) is generated.
- ✓ When you save the data these requests are passed to the database in the order they were generated.
- ✓ So, multiple people can work in the same table at the same time with no worries
 - Risk of data loss is low
 - Still possible for someone else to change data after you change it.
 - Versioning is an option to record the changes (DB Admins are owner of these things in production)

SQL Command Categories



Data Query Language

Select

Data Definition Language

Create Table

Alter Table

Drop Table

Data Manipulation Language

Insert

Update

Delete

Data Control Language

Grant

Revoke

Data Types in SQL


Defines the type of data the column holds

	sid_Department	dept_no	dept_name
1	1	d001	Marketing
2	2	d002	Finance
3	3	d003	Human Resources
4	4	d004	Production
5	5	d005	Development
6	6	d006	Quality Management
7	7	d007	Sales
8	8	d008	Research
9	9	d009	Customer Service
10	10	d010	Rocketry & Telemetry
11	11	d011	Data Science & Analytics

Integer

Character

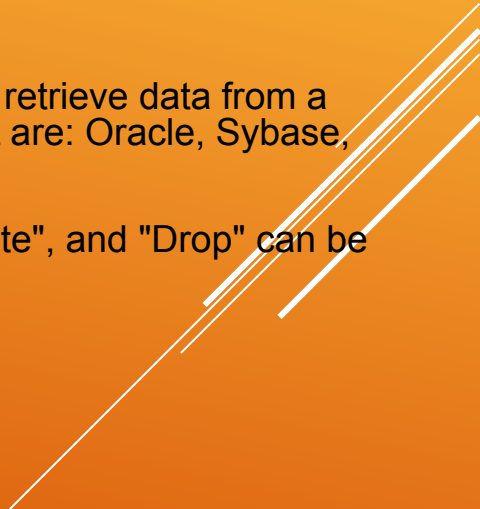
Different Data Types



Introduction



What is SQL?

- ✓ SQL (pronounced "ess-que-el") stands for Structured Query Language
 - ✓ SQL is used to communicate with a database.
 - ✓ It is the standard language for relational database management systems
 - ✓ SQL statements are used to perform tasks such as update data on a database, or retrieve data from a database. Some common relational database management systems that use SQL are: Oracle, Sybase, Microsoft SQL Server, Access, Ingres, etc.
 - ✓ The standard SQL commands such as "Select", "Insert", "Update", "Delete", "Create", and "Drop" can be used to accomplish almost everything that one needs to do with a database.
 - ✓ SQL programming can be used to perform multiple actions on data such as :
 - Querying
 - Inserting
 - Updating
 - Deleting
 - Extracting etc.
- 

Primary & Foreign Key

A Primary key is used to ensure data in the specific column is Unique and Not Null

Primary key

=

Unique

+

Not Null

A foreign key is a column or group of columns in a relational database table that provides a link between data in two tables. It is a column (or columns) that references a column (most often the primary key) of another table

Primary & Foreign Key - Contd..

□ Primary key:


- no null values
- unique identification

□ Foreign key:

- correspond to the values of the primary key in another table

```
create table dep (  
  depid int not null,  
  depname varchar(100),  
  depaddress varchar(255),  
  PRIMARY KEY (depid))
```

```
create table emp (  
  empid int,  
  empname varchar(100),  
  empadd varchar(255),  
  depid int,  
  PRIMARY KEY (empid),  
  FOREIGN KEY (depid) REFERENCES dep(depid))
```



SQL – First Step




- ❑ Create a Database
- ❑ Create a Table
- ❑ Insert Table Data
- ❑ Update Table Data
- ❑ Alter Database
- ❑ Alter Table
- ❑ Alter Column
- ❑ Duplicate a Database
- ❑ Drop a Database
- ❑ Database Backup
- ❑ Restore a Database
- ❑ Rename Database



TABLE BASICS

A relational database system contains one or more objects called tables. The data or information for the database are stored in these tables. Tables are uniquely identified by their names and are comprised of columns and rows. Columns contain the column name, data type, and any other attributes for the column. Rows contain the records or data for the columns. Here is a sample table called "employee".

Employee Name	Employee Id	Manager Name	Division
Satyajit Pattnaik	901	Rakesh Dash	1
Ramesh Sahoo	902	Rakesh Dash	1
Rakesh Dash	903	Subrat Pal	1
Subrat Pal	904	Santosh Das	1
Santosh Das	905	-	1



The LIKE pattern matching operator can also be used in the conditional selection of the where clause. Like is a very powerful operator that allows you to select only rows that are "like" what you specify. The percent sign "%" can be used as a wildcard to match any possible character that might appear before or after the characters specified.

This SQL statement will match any first names that start with 'Er'. Strings must be in single quotes. Or you can specify:

For example:

```
select first, last, city
from empinfo
where first LIKE 'Er%';
```

This statement will match any last names that end in a 's'.

```
select first, last
from empinfo
where last LIKE '%s';
```


This will only select rows where the first name equals 'Eric' exactly

```
select * from empinfo
where first = 'Eric';
```



SELECTING THE DATA

Sample Table: empinfo					
first	last	id	age	city	state
John	Jones	99980	45	Payson	Arizona
Mary	Jones	99982	25	Payson	Arizona
Eric	Edwards	88232	32	San Diego	California
Mary Ann	Edwards	88233	32	Phoenix	Arizona
Ginger	Howell	98002	42	Cottonwood	Arizona
Sebastian	Smith	92001	23	Gila Bend	Arizona
Gus	Gray	22322	35	Bagdad	Arizona
Mary Ann	May	32326	52	Tucson	Arizona
Erica	Williams	32327	60	Show Low	Arizona
Leroy	Brown	32380	22	Pinetop	Arizona
Elroy	Cleaver	32382	22	Globe	Arizona



Enter the following sample select statements in the SQL Interpreter Form at the bottom of this page. Before you press "submit", write down your expected results. Press "submit", and compare the results.

```
select first, last, city from empinfo;
```

```
select last, city, age from empinfo  
where age > 30;
```

```
select first, last, city, state from empinfo  
where first LIKE 'J%';
```

```
select * from empinfo;
```

```
select first, last, from empinfo  
where last LIKE '%s';
```

```
select first, last, age from empinfo  
where last LIKE '%illia%';
```

```
select * from empinfo where first = 'Eric';
```



CREATING TABLES

<p>The create table statement is used to create a new table. Here is the format of a simple create table statement:</p>	<pre>create table "tablename" ("column1" "data type", "column2" "data type", "column3" "data type");</pre>
<p>Format of create table if you were to use optional constraints:</p>	<pre>create table "tablename" ("column1" "data type" [constraint], "column2" "data type" [constraint], "column3" "data type" [constraint]); [] = optional</pre>
<p>Note: You may have as many columns as you'd like, and the constraints are optional.</p>	<p>Example:</p> <pre>create table employee (first varchar(15), last varchar(20), age number(3), address varchar(30), city varchar(20), state varchar(20));</pre>



To create a new table,

- ❑ enter the keywords create table followed by the table name,
- ❑ followed by an open parenthesis,
- ❑ followed by the first column name,
- ❑ followed by the data type for that column,
- ❑ followed by any optional constraints, and followed by a closing parenthesis.

It is important to make sure you use an open parenthesis before the beginning table, and a closing parenthesis after the end of the last column definition.

Make sure you separate each column definition with a comma. **All SQL statements should end with a ";"**.





The table and column names;


- must start with a letter
- and can be followed by letters, numbers, or underscores
- not to exceed a total of 30 characters in length.

Do not use any SQL reserved keywords as names for tables or column names (such as "select", "create", "insert", etc).


Data types specify what the type of data can be for that particular column. If a column called "Last_Name", is to be used to hold names, then that particular column should have a **"varchar"** (variable-length character) data type.



Here are the most common Data types:



<code>char(size)</code>	Fixed-length character string. Size is specified in parenthesis. Max 255 bytes.
<code>varchar(size)</code>	Variable-length character string. Max size is specified in parenthesis.
<code>number(size)</code>	Number value with a max number of column digits specified in parenthesis.
<code>date</code>	Date value
<code>number(size, d)</code>	Number value with a maximum number of digits of "size" total, with a maximum number of "d" digits to the right of the decimal.



What are constraints?

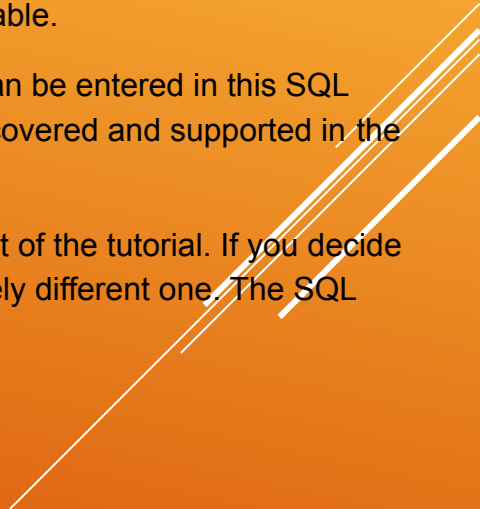


When tables are created, it is common for one or more columns to have constraints associated with them. A constraint is basically a rule associated with a column that the data entered into that column must follow.


- For example, a "unique" constraint specifies that no two records can have the same value in a particular column. They must all be unique.
- The other two most popular constraints are "not null" which specifies that a column can't be left blank, and "primary key". A "primary key" constraint defines a unique identification of each record (or row) in a table.

All of these and more will be covered in the future Advanced release of this Tutorial. Constraints can be entered in this SQL interpreter, however, they are not supported in this Intro to SQL tutorial & interpreter. They will be covered and supported in the future release of the Advanced SQL tutorial - that is, if "response" is good.

It's now time for you to design and create your own table. You will use this table throughout the rest of the tutorial. If you decide to change or redesign the table, you can either drop it and recreate it or you can create a completely different one. The SQL statement drop will be covered later.



Create Table Exercise



You have just started a new company. It is time to hire some employees. You will need to create a table that will contain the following information about your new employees:

firstname, last name, title, age, and salary.

After you create the table, you should receive a small form on the screen with the appropriate column names. If you are missing any columns, you need to double check your SQL statement and recreate the table. Once it's created successfully, go to the "Insert" lesson.

IMPORTANT: When selecting a table name, it is important to select a unique name that no one else will use or guess. Your table names should have an underscore followed by your initials and the digits of your birth day and month.

For example, Tom Smith, who was born on November 2nd, would name his table myemployees_ts0211 Use this convention for all of the tables you create. Your tables will remain on a shared database until you drop them, or they will be cleaned up if they aren't accessed in 4-5 days

INSERT

The insert statement is used to insert or add a row of data into the table.

To insert records into a table, enter the key words insert into followed by the table name, followed by an open parenthesis, followed by a list of column names separated by commas, followed by a closing parenthesis, followed by the keyword values, followed by the list of values enclosed in parenthesis. The values that you enter will be held in the rows and they will match up with the column names that you specify.

Strings should be enclosed in single quotes, and numbers should not.

```
insert into "tablename"  
(first_column,...last_column)  
values (first_value,...last_value);
```

Example:

```
insert into empinfo  
(first, last, id, age, city, state)  
values ('Luke', 'Duke', 45454,  
'22', 'Hazard Co', 'Georgia');
```

Note: All strings should be enclosed between single quotes: 'string'

Insert statement exercise

A horizontal bar composed of seven colored segments: blue, light blue, green, orange, pink, purple, and light blue. Each segment contains a small white circle.

It is time to insert data into your new employee table.

Your first three employees are the following:

Jonie Weber, Secretary, 28, 19500.00

Potsy Weber, Programmer, 32, 45300.00

Dirk Smith, Programmer II, 45, 75020.00

Enter these employees into your table first, and then insert at least 5 more of your own list of employees in the table.

After they're inserted into the table, enter select statements to:

Select all columns for everyone in your employee table.

Select all columns for everyone with a salary over 30000.

Select first and last names for everyone that's under 30 years old.

Select first name, last name, and salary for anyone with "Programmer" in their title.

Select all columns for everyone whose last name contains "ebe".

Select the first name for everyone whose first name equals "Potsy".

Select all columns for everyone over 80 years old.

Select all columns for everyone whose last name ends in "ith".

Create at least 5 of your own select statements based on specific information that you'd like to retrieve.

Three parallel white diagonal lines on an orange background, extending from the bottom right towards the top right.

UPDATE

The update statement is used to update or change records that match a specified criteria. This is accomplished by carefully constructing a where clause.

```
update "tablename"  
set "columnname" =  
    "newvalue"  
[, "nextcolumn" =  
    "newvalue2"...]  
where "columnname"  
    OPERATOR "value"  
[and|or "column"  
    OPERATOR "value"];  
[] = optional
```

Examples:

```
update phone_book  
set area_code = 623  
where prefix = 979;
```

```
update phone_book  
set last_name = 'Smith',  
    prefix=555, suffix=9292  
where last_name = 'Jones';
```

```
update employee  
set age = age+1  
where first_name='Mary' and  
    last_name='Williams';
```

Update statement exercises

After each update, issue a select statement to verify your changes.

1. Jonie Weber just got married to Bob Williams. She has requested that her last name be updated to Weber-Williams.
2. Dirk Smith's birthday is today, add 1 to his age.
3. All secretaries are now called "Administrative Assistant". Update all titles accordingly.
4. Everyone that's making under 30000 are to receive a 3500 a year raise.
5. Everyone that's making over 33500 are to receive a 4500 a year raise.
6. All "Programmer II" titles are now promoted to "Programmer III".
7. All "Programmer" titles are now promoted to "Programmer II".

Create at least 5 of your own update statements and submit them.

DELETING RECORDS

The delete statement is used to delete records or rows from the table.

Examples:

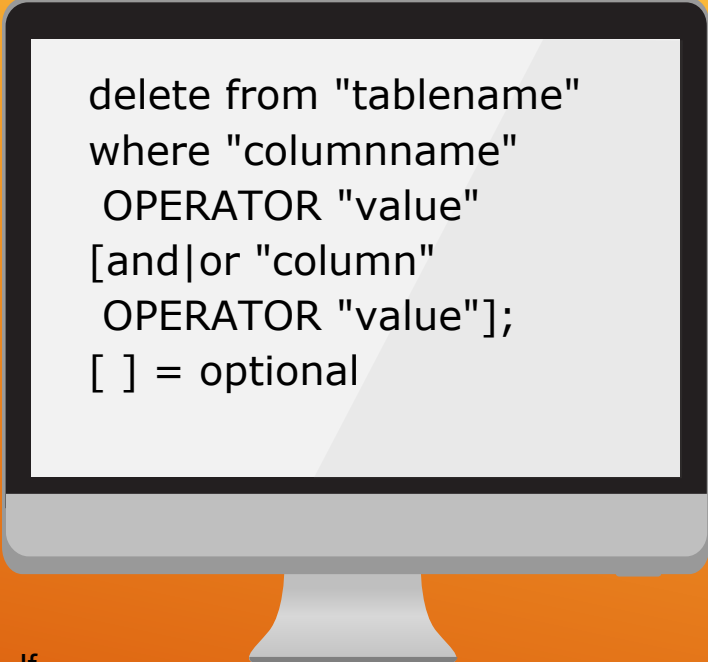
```
delete from employee;
```

Note: if you leave off the where clause, **all records will be deleted!**

```
delete from employee  
where lastname = 'May';
```

```
delete from employee  
where firstname = 'Mike' or firstname = 'Eric';
```

To delete an entire record/row from a table, enter "delete from" followed by the table name, followed by the where clause which contains the conditions to delete. If you leave off the where clause, all records will be deleted.



```
delete from "tablename"  
where "columnname"  
    OPERATOR "value"  
[and|or "column"  
    OPERATOR "value"];  
[ ] = optional
```

DROP A TABLE

- The **drop table** command is used to delete a table and all rows in the table.
- To delete an entire table including all of its rows, issue the **drop table** command followed by the tablename.
- **drop table** is different from deleting all of the records in the table.
- Deleting all of the records in the table leaves the table including column and constraint information.
- Dropping the table removes the table definition as well as all of its rows.




```
drop table "tablename"
```

Example:

```
drop table  
myemployees_ts0211;
```

JOINS


- In order to sort a table, we need to use the **ORDER BY** Clause
- In order to get the results sorted based on certain columns, we need to use this. Usage of **asc** or **desc** has to be defined to have the results in ascending or descending orders. Default value being **asc** (ascending)



```
select field1, field2 from  
tablename where <clause>  
order by field1
```


What are constraints?

- In most of the real world problems, we might need data from multiple tables, that's where **Joins** comes into picture.



JOINS - Right Join

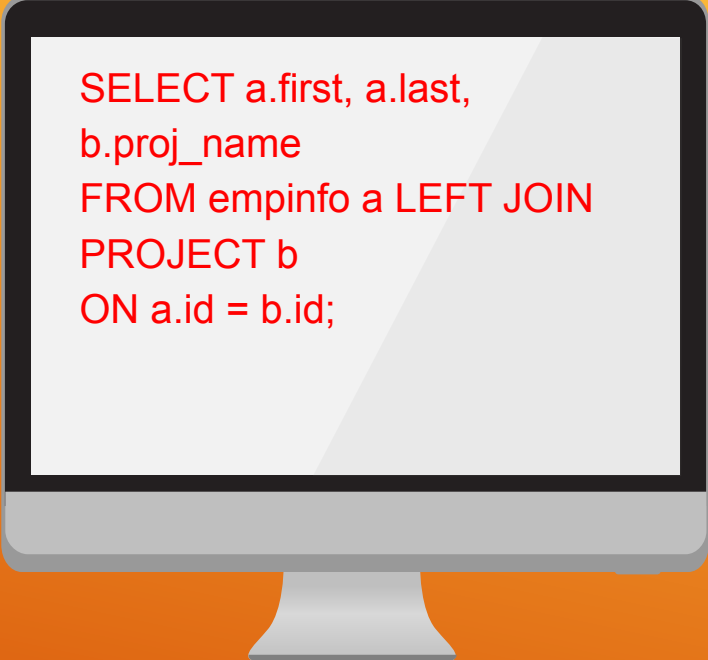
- The RIGHT JOIN keyword returns all records from the right table (table2), and the matched records from the left table (table1).



```
SELECT a.first, a.last,  
b.proj_name  
FROM empinfo a RIGHT  
JOIN PROJECT b  
ON a.id = b.id;
```

JOINS - Left Join

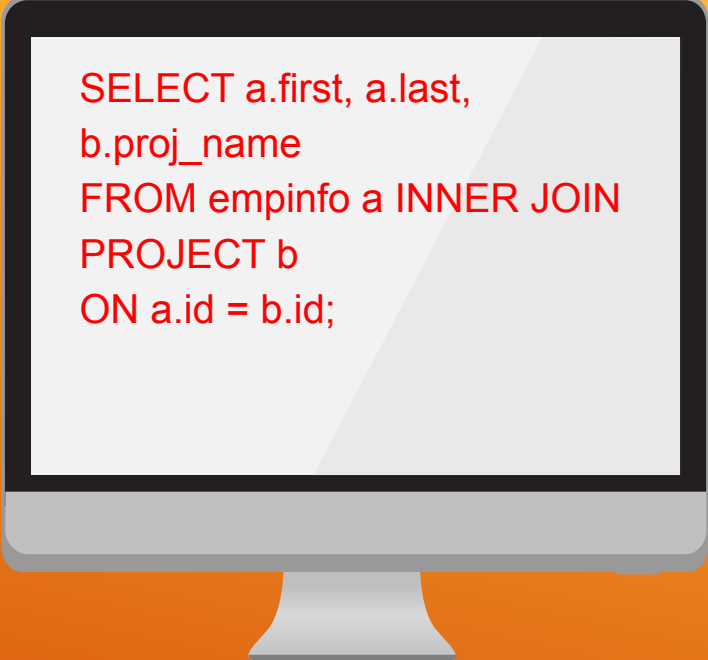
- The LEFT JOIN keyword returns all records from the left table (table1), and the matched records from the right table (table2). The result is NULL from the right side, if there is no match.



```
SELECT a.first, a.last,  
       b.proj_name  
FROM empinfo a LEFT JOIN  
PROJECT b  
ON a.id = b.id;
```

JOINS - Inner Join

- The INNER Join keyword selects records that have matching values in both tables




```
SELECT a.first, a.last,  
b.proj_name  
FROM empinfo a INNER JOIN  
PROJECT b  
ON a.id = b.id;
```

Alter Table

ALTER TABLE
tablename
ADD columnname
datatype;

ALTER TABLE
tablename
DROP COLUMN
columnname ;

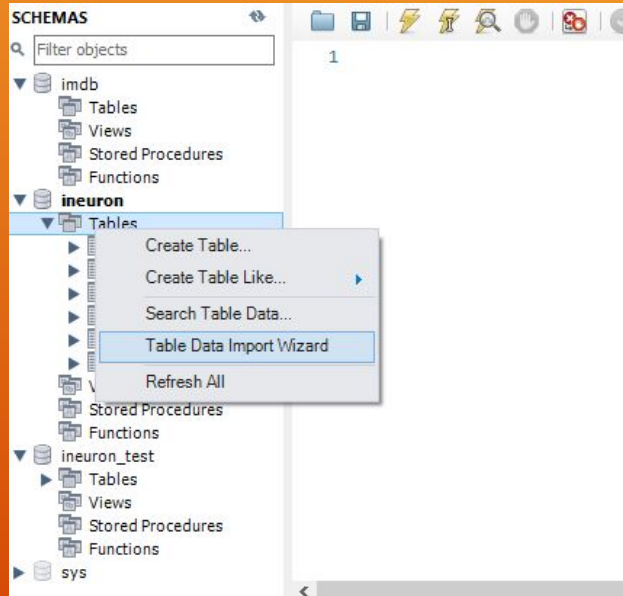


ALTER TABLE
CUST_DETAILS
ADD AGE INT;

ALTER TABLE
CUST_DETAILS
DROP COLUMN AGE;

Data Importing

1. Manual Importing



Data Importing

1. Importing through command line

- Open MySQL Workbench, Create a new database to store the tables you'll import (eg- FacilityServices) → Then create the table using CREATE query
- Copy the MySQL bin directory path: C:\Program Files\MySQL\MySQL Server 8.0\bin
- Go to the folder in command line by using: cd path
- Connect to MySQL database: mysql -u root -p (root is basically your username)
- If you are logged in successfully, then set the global variables by using below command so that the data can be imported from local computer folder.
 - mysql> SET GLOBAL local_infile = 1;
 - Query OK, 0 rows affected (0.00 sec)
 - (you've just instructed MySQL server to allow local file upload from your computer)
- Quit current server connection (mysql> quit)
- Load the file from CSV file to the MySQL database. In order to do this, please follow the commands: (We'll connect with the MySQL server again with the local-infile system variable. This basically means you want to upload data into a file from a local machine)
 - mysql --local-infile=1 -u root -p (give password)
 - Show Databases; (It'll show all the databases in MySQL server.)
 - mysql> USE dbase; (makes the database that you had created in step 1 as default schema to use for the next sql scripts)

```
mysql> LOAD DATA LOCAL INFILE 'fullpath\file.csv'
      INTO TABLE tablename
      FIELDS TERMINATED BY ','
      ENCLOSED BY '"'
      LINES TERMINATED BY '\r\n' IGNORE 1 ROWS;
```

- **Note: VERY IMP** - Please replace single backward (\) slash in the path with double back slashes (\\) instead of single slash

Data Exporting

Server → Data Export

The screenshot shows the MySQL Data Export wizard for a local instance of MySQL 8.0. The interface is divided into several sections:

- Object Selection:** Contains two panes. The left pane, titled "Tables to Export", shows a list of databases: "imdb", "ineuron", "ineuron_test", and "sys". The "imdb" database is selected. The right pane, titled "Schema Objects", shows a list of objects: "genre", "movie", and "person". All three are selected.
- Export Options:** Includes a "Refresh" button and a status "3 tables selected". Below this, there are checkboxes for "Dump Stored Procedures and Functions", "Dump Events", and "Dump Triggers", all of which are currently unchecked.
- Export Options:** Includes a radio button for "Export to Dump Project Folder" (selected) and a text field showing the path "C:\Users\User\Documents\dumps\Dump20210912". Below this, there is a radio button for "Export to Self-Contained File" and a text field showing the path "C:\Users\User\Documents\dumps\Dump20210912.sql".
- Advanced Options:** A button labeled "Advanced Options..." is located in the top right corner.
- Start Export:** A button labeled "Start Export" is located at the bottom right.

At the bottom left, there is a message: "Press [Start Export] to start..."

Aggregate Functions



- ✓ Sometimes we examine & analyse data of varying magnitudes, hence we realise the need of grouping similar types of values together & look them at as one bunch.

For example: Consider a table containing data consisting of the marks scored by students in their 12th board exams. While you would want to know how the students performed in all the subjects put together, it is equally important to see how they performed in each subject. You can gain even further insights if you group these students by state. Hence, it is imperative that you learn the usage of aggregate functions in your queries.

- ✓ `groupby()` → To aggregate values of a column C1 'grouped by' a certain column C2.
- ✓ `count()` → Count the number of rows.
- ✓ `min()`, `max()` → Finding the minimum & maximum values for a particular column

For example: `select min(score) from employee;`


- ✓ `avg()` → Find the average

ORDERING




Quite often, you would want to display the retrieved records in a particular order, for example, in increasing order of income, joining date, alphabetical order, etc. This is commonly useful when you are making a report or presenting the data to someone else.

```
select firstName from  
employees order by  
firstName asc limit 3;
```



HAVING CLAUSE



You have already learnt how to filter individual values based on a given condition.
But how do you do this on grouped values?

Suppose your manager asks you to count all the employees whose salaries are more than the average salary in that particular department.

Now, intuitively, you know that two aggregate functions would be used here, namely, `count()` and `avg()`. You decide to apply the 'where' condition on the average salary of the department, but to your surprise, the query fails. This is exactly what the having clause is for.

The 'having' clause is typically used when you have to apply a filter condition on an 'aggregated value'. This is because the 'where' clause is applied before aggregation takes place, and thus, it is not useful when you want to apply a filter on an aggregated value.

The **HAVING** clause was added to SQL because the **WHERE** keyword cannot be used with aggregate functions.

Exercise

Keywords in SQL

List the keywords below in order of their occurrence in a query:

1. group by
2. order by
3. select
4. where
5. from
6. limit
7. having

Answer is:

STRING FUNCTIONS

Used to manipulate the string data and make it more understandable for analysis.

For example: amitabhbachchan, or Amitabh Bachchan, which one of them is more readable, obviously the later one right.

concat → Concatenates two strings

substr/substring → Extracts a substring from a string

upper → Converts a string to uppercase

lower → Converts a string to lowercase

character_length → Calculates the length of the string variable.

mid → find the middle elements → MID(str,pos,len)

https://www.w3schools.com/sql/sql_ref_mysql.asp

DATE & TIME FUNCTIONS

Used to manipulate the date & time columns

For example: If you want to change the date format, or just want to see the exact day, or so on.

Let's create a transaction_details table and play around.

datediff → Return the number of days between the two date values

```
SELECT DATEDIFF(sysdate(), order_date) from transaction_details
```

date_format → Format a date variable

```
SELECT DATE_FORMAT("2017-06-15", "%Y");
```

day → Return the day of the month for a date.

```
SELECT DAY("2017-06-15");
```

quarter → Return the quarter of the year for a date.

```
SELECT QUARTER("2017-06-15");
```

addday → Add days to the date variable

```
SELECT ADDDATE("2017-06-15", INTERVAL 10 DAY);
```

and so on....

https://www.w3schools.com/sql/func_mysql_adddate.asp

REGEX

A decorative horizontal bar consisting of seven colored segments: blue, cyan, green, orange, pink, purple, and light blue. Each segment contains a small white circle.

So far you have already known about the wildcards like “like” operator, but in cases wildcards may fall short for some advanced use cases, regular expressions comes into picture.

Regex, or Regular expressions, is a sequence of characters, used to search and locate specific sequences of characters that match a pattern.

Example 1: Match beginning of the name

```
select * from customer WHERE email REGEXP '^w'
```

A series of three parallel white diagonal lines on an orange background, extending from the bottom right towards the top right.

REGEX

0

0

0

0

0

0

Example 2: Find the customers, which email address having a 'z', 'v' or 'p'

```
SELECT * FROM customer  
WHERE email REGEXP "[zvp]";
```

Example 3: Find the customers, which email address containing characters from 'x' to 'z'


```
SELECT * FROM customer  
WHERE email REGEXP "[x-z]";
```

NESTED QUERIES

By now, you know that a database is a collection of multiple related tables. Now, while generating insights from data, you may need to refer to these multiple tables in a query. There are two ways to deal with such types of queries. These include the following:

1. Joins
2. Nested queries/Subqueries

A subquery is called an inner query while the query that contains the subquery is called an outer query. A subquery can be used anywhere that expression is used and must be closed in parentheses.




```
SELECT lastName,firstName
FROM employees
WHERE
    office_Code IN (SELECT
        office_Code
        FROM
            offices
        WHERE
            country = 'USA');
```


STORED PROCEDURES

A stored procedure is a prepared SQL code that you can save, so the code can be reused over and over again.

So if you have an SQL query that you write over and over again, save it as a stored procedure, and then just call it to execute it.

You can also pass parameters to a stored procedure, so that the stored procedure can act based on the parameter value(s) that is passed.

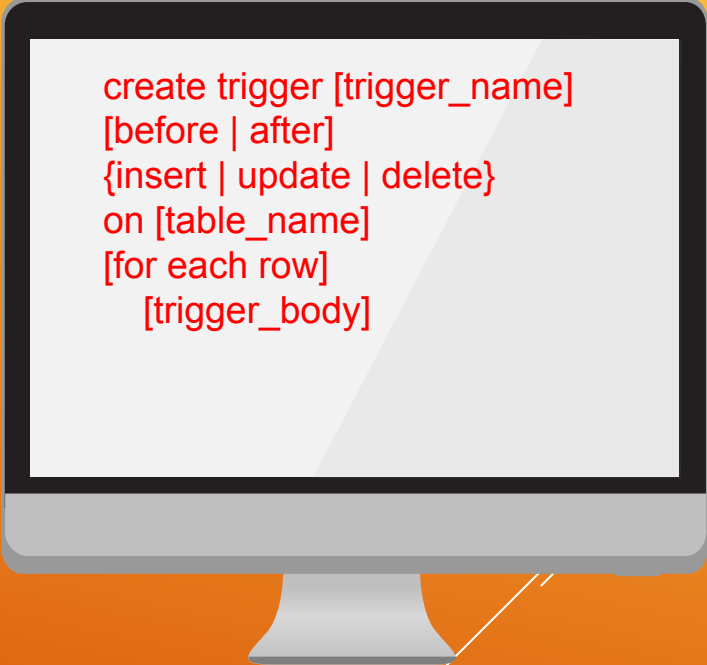


```
CREATE PROCEDURE  
SelectAllCustomers  
AS  
SELECT * FROM  
Customers  
GO;
```

TRIGGERS

A trigger is a stored procedure in database which automatically invokes whenever a special event in the database occurs. For example, a trigger can be invoked when a row is inserted into a specified table or when certain table columns are being updated.

- DDL Trigger
- DML Trigger
- Logon Trigger



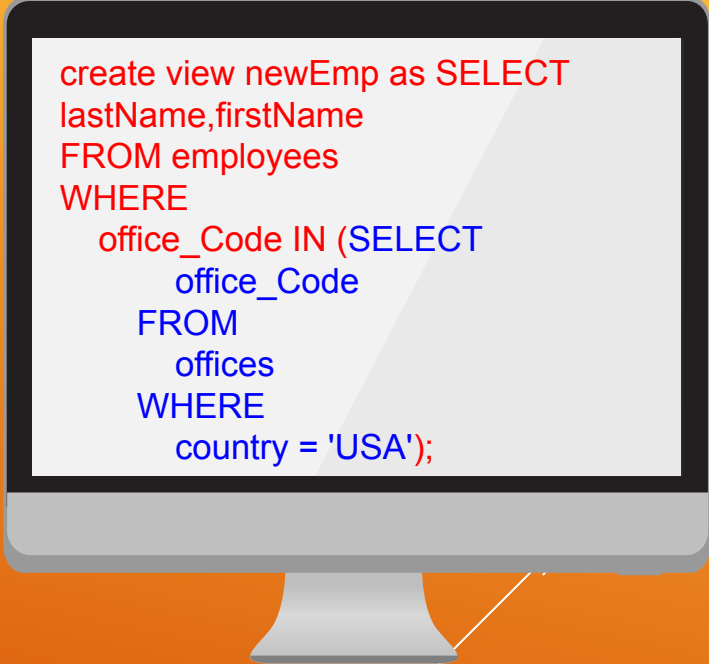
```
create trigger [trigger_name]
[before | after]
{insert | update | delete}
on [table_name]
[for each row]
[trigger_body]
```

VIEWS

Views are virtual tables that do not store any data of their own but display data stored in other tables.

Advantages:

- Hide the complexity of data.
- Act as aggregated tables.
- If you are doing an user level access control, you can give an user access to a view without giving access to the tables behind it
- It can allow for massive performance improvements.

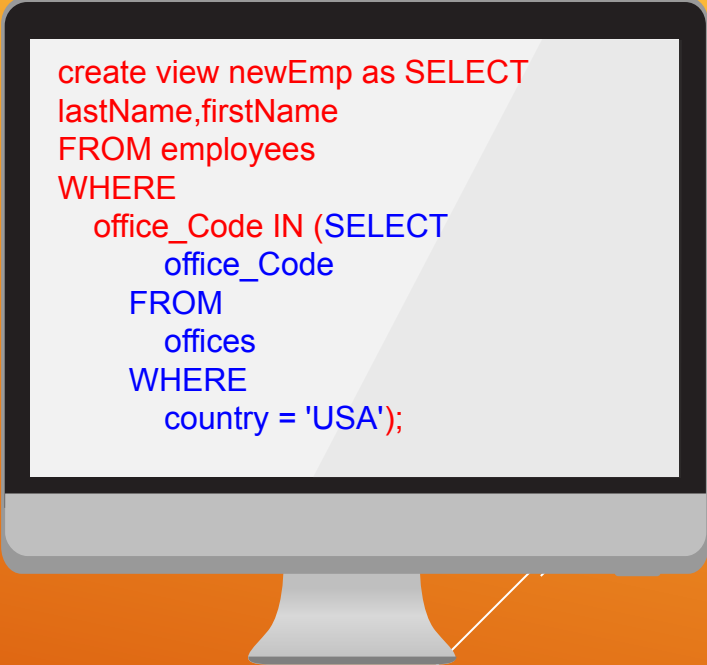


```
create view newEmp as SELECT  
lastName,firstName  
FROM employees  
WHERE  
    office_Code IN (SELECT  
        office_Code  
    FROM  
        offices  
    WHERE  
        country = 'USA');
```

VIEWS

Disadvantages:


- If done wrong, it can result in performance issues.
- You may not be able to update the view, forcing you back to the original tables.



```
create view newEmp as SELECT  
lastName,firstName  
FROM employees  
WHERE  
    office_Code IN (SELECT  
        office_Code  
    FROM  
        offices  
    WHERE  
        country = 'USA');
```

VIEWS

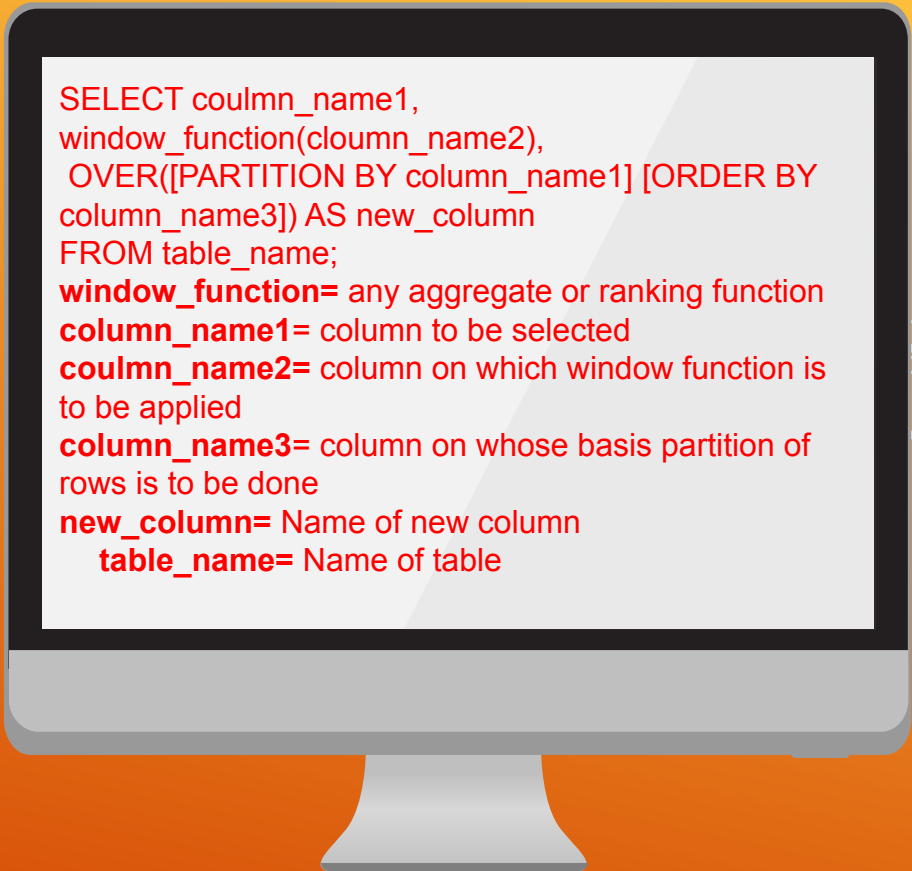
Anatomy of a View



WINDOW FUNCTION

Window functions apply aggregate and ranking functions over a particular window (set of rows). OVER clause is used with window functions to define that window. OVER clause does two things :

- Partitions rows into form set of rows. (PARTITION BY clause is used)
- Orders rows within those partitions into a particular order. (ORDER BY clause is used)




```
SELECT coulumn_name1,  
window_function(cloumn_name2),  
  OVER([PARTITION BY column_name1] [ORDER BY  
column_name3]) AS new_column  
FROM table_name;
```

window_function= any aggregate or ranking function
column_name1= column to be selected
coulumn_name2= column on which window function is to be applied
column_name3= column on whose basis partition of rows is to be done
new_column= Name of new column
table_name= Name of table

WINDOW FUNCTION



- Row_number() → Gives a sequential integer to every row within its partition
 - Rank() functions → Ranking records
 - First_value() functions → Returns the value of the specific expression with respect to first row in the window frame
- 

SQL-Python Connectivity



- pymysql → This package contains a pure-Python MySQL client library, based on PEP 249
→ pip install pymysql

```
dbcon = pymysql.connect(host="localhost", user="root", password="root", database="breast_cancer")
```



SUMMARY



You learnt about the crux of SQL queries. These are the basic to intermediate level concepts of SQL that will be tested rigorously in your interviews. Hence, practice these points again and again to gain a deep understanding of the several SQL statements.

Leave your top three takeaways from this session in the comments section below.



Exercise



Write a query to retrieve the names of all employees who have an age greater than what Gus Gray has.

Answer is:




Exercise

Write a query to retrieve the names of all employees who have an age greater than what Gus Gray has, and store the output in a view.

Answer is:

Exercise



Write a query to retrieve the names of all employees who have an age greater than what Gus Gray has, and store the output in a view.

Answer is:

